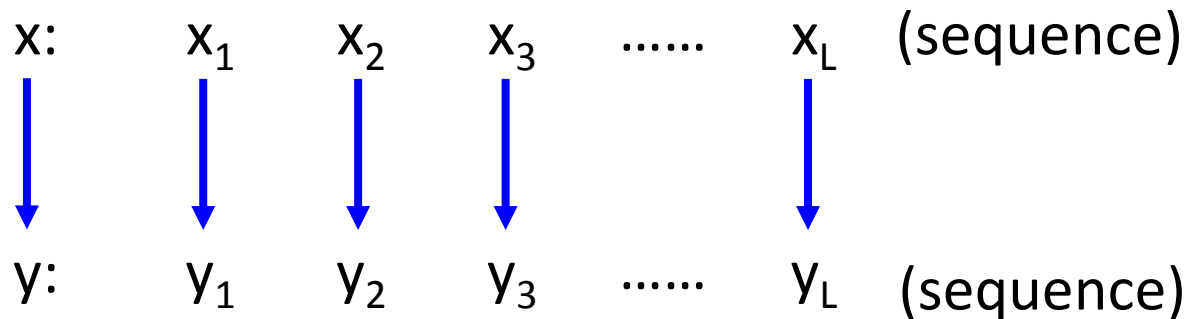


# Sequence Labeling Problem

# Sequence Labeling

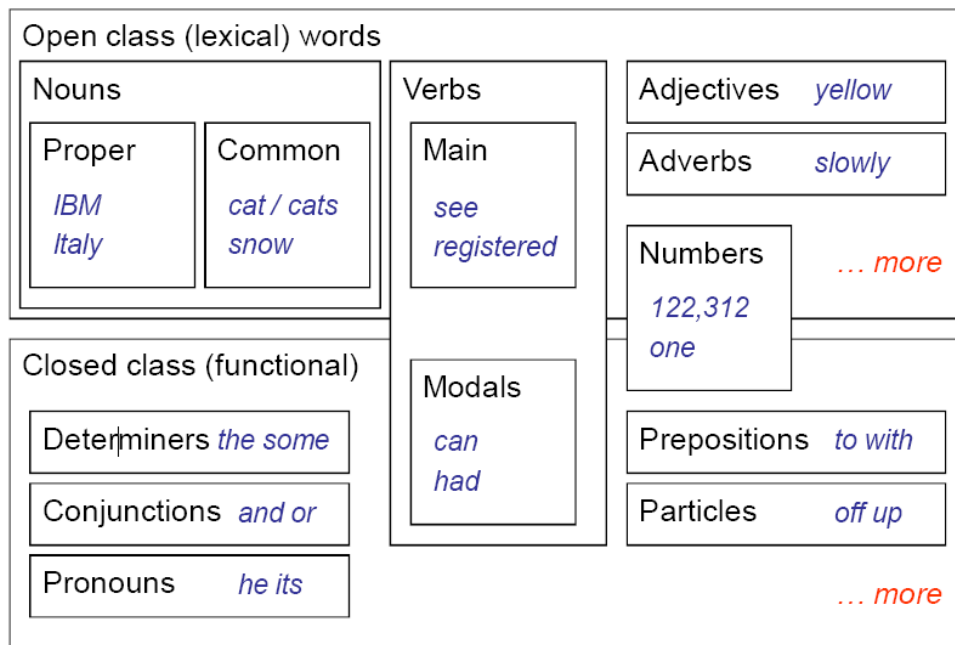
$$f : \underset{\text{Sequence}}{X} \rightarrow \underset{\text{Sequence}}{Y}$$



RNN can handle this task, but there are other methods based on structured learning (two steps, three problems).

# Example Task

- POS tagging
  - Annotate each word in a sentence with a part-of-speech.



John saw the saw.

↓ ↓ ↓ ↓

PN V D N

- Useful for subsequent syntactic parsing and word sense disambiguation, etc.

# Example Task

- POS tagging

John saw the saw.  
↓ ↓ ↓ ↓  
PN V D N

The problem cannot be solved  
without considering the sequences.

- “saw” is more likely to be a verb V rather than a noun N
- However, the second “saw” is a noun N because a noun N is more likely to follow a determiner.

# Outline

Hidden Markov Model (HMM)



Conditional Random Field (CRF)



Structured Perceptron/SVM



Towards Deep Learning

# Outline

Hidden Markov Model (HMM)



Conditional Random Field (CRF)



Structured Perceptron/SVM



Towards Deep Learning

# HMM

- How you generate a sentence?

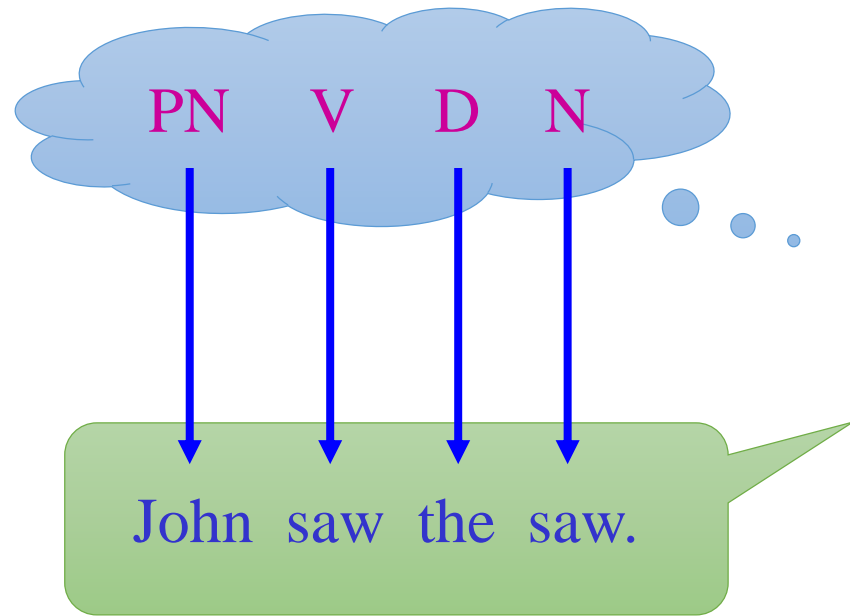
Just the assumption  
of HMM

## Step 1

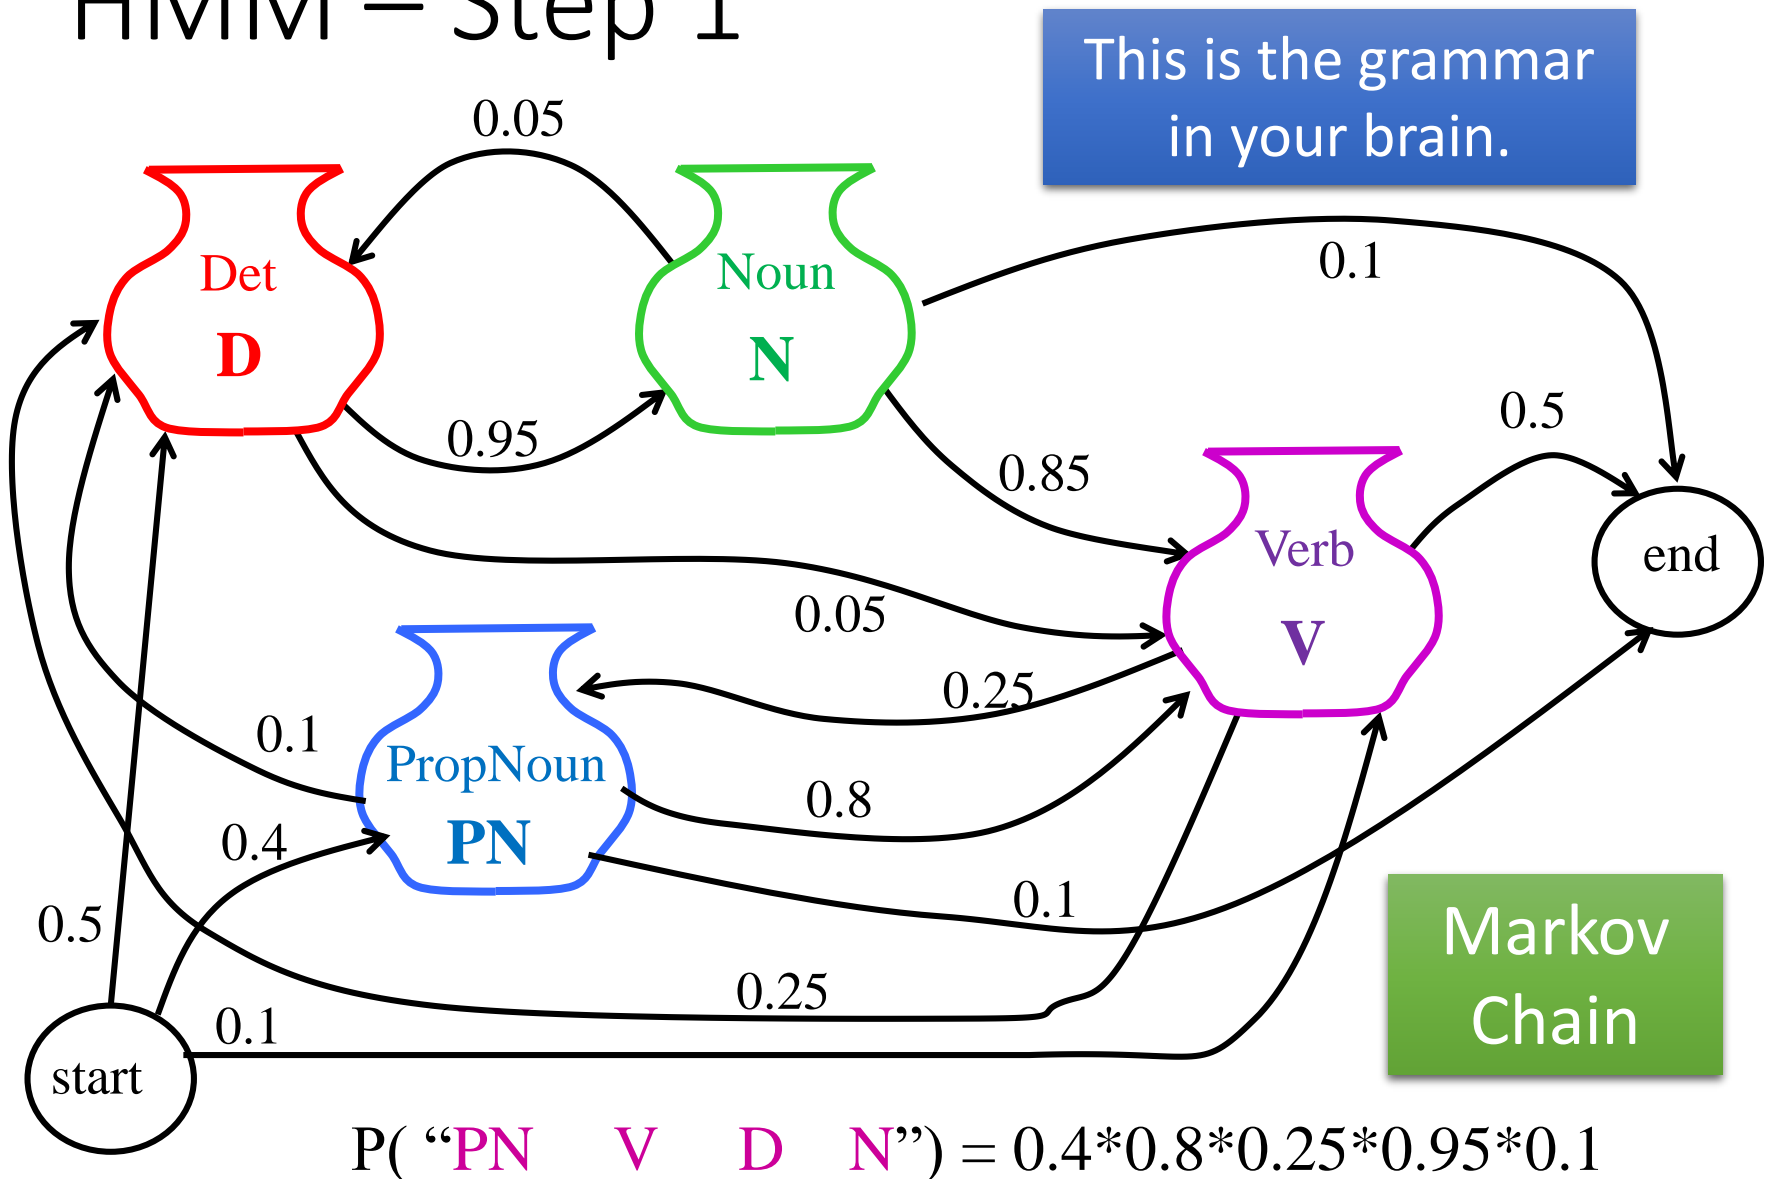
- Generate a POS sequence
- Based on the grammar

## Step 2

- Generate a sentence based on the POS sequence
- Based on a dictionary

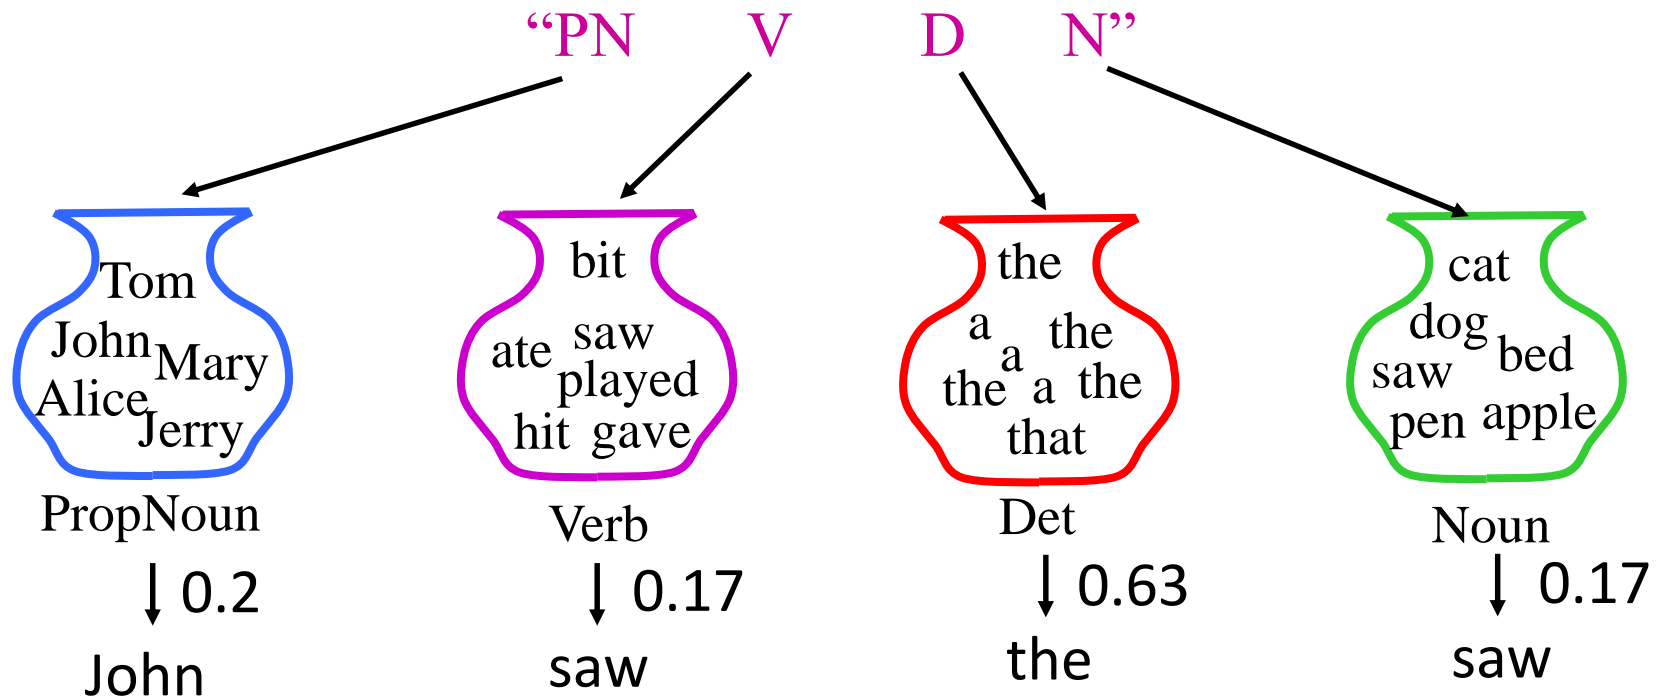


# HMM – Step 1



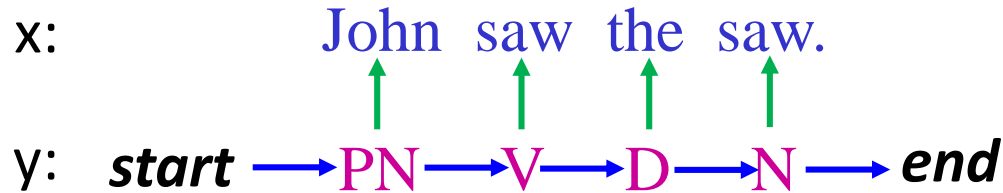


# HMM – Step 2



$$P(\text{"John saw the saw"} \mid \text{"PN V D N"}) \\ = 0.2 * 0.17 * 0.63 * 0.17$$

# HMM



$$P(x,y) = P(y)P(x|y)$$

$$\begin{aligned} P(y) &= P(PN|start) \\ &\times P(V|PN) \\ &\times P(D|V) \\ &\times P(N|D) \end{aligned}$$

$$\begin{aligned} P(x|y) &= P(John|PN) \\ &\times P(saw|V) \\ &\times P(the|D) \\ &\times P(saw|N) \end{aligned}$$

# HMM

x: John saw the saw.

$$x = x_1, x_2 \cdots x_L$$

y: PN V D N

$$y = y_1, y_2 \cdots y_L$$

$$P(x, y) = P(y)P(x|y)$$

## Step 1

$$P(y) = P(y_1 | start) \times \prod_{l=1}^{L-1} P(y_{l+1} | y_l) \times P(end | y_L)$$

Transition probability

## Step 2

$$P(x|y) = \prod_{l=1}^L P(x_l | y_l)$$

Emission probability

# HMM

## – Estimating the probabilities

- How can I know  $P(V|PN)$ ,  $P(\text{saw}|V)$  ..... ?
- Obtaining from training data

### Training Data:

- $(x^1, \hat{y}^1)$  1 Pierre/NNP Vinken/NNP ,/, 61/CD years/NNS old/JJ ,/, will/MD join/VB the/DT board/NN as/IN a/DT nonexecutive/JJ director/NN Nov./NNP 29/CD ./.
- $(x^2, \hat{y}^2)$  2 Mr./NNP Vinken/NNP is/VBZ chairman/NN of/IN Elsevier/NNP N.V./NNP ,/, the/DT Dutch/NNP publishing/VBG group/NN ./.
- $(x^3, \hat{y}^3)$  3 Rudolph/NNP Agnew/NNP ,/, 55/CD years/NNS old/JJ and/CC chairman/NN of/IN Consolidated/NNP Gold/NNP Fields/NNP PLC/NNP ,/, was/VBD named/VBN a/DT nonexecutive/JJ director/NN of/IN this/DT British/JJ industrial/JJ conglomerate/NN ./.

⋮

# HMM

## – Estimating the probabilities

$$P(x, y) = \underbrace{P(y_1 | start)} \prod_{l=1}^{L-1} \underbrace{P(y_{l+1} | y_l)} \underbrace{P(end | y_L)} \prod_{l=1}^L \underbrace{P(x_l | y_l)}$$

$$\underbrace{P(y_{l+1} = s' | y_l = s)} = \frac{count(s \rightarrow s')}{count(s)}$$

(s and s' are tags)

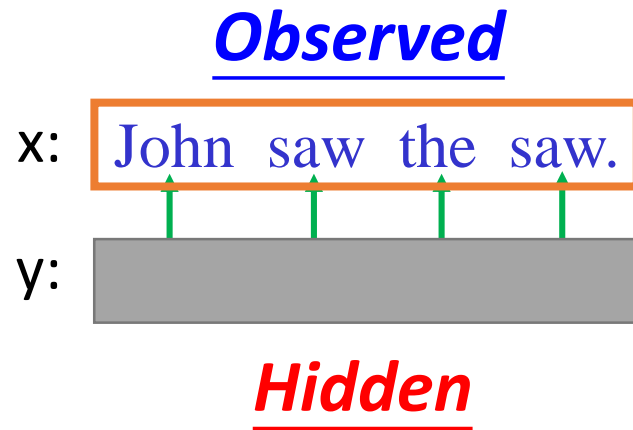
$$\underbrace{P(x_l = t | y_l = s)} = \frac{count(s \rightarrow t)}{count(s)}$$

(s is tag, and t is word)

So simple 😊

# HMM – How to do POS Tagging?

- We can compute  $P(x, y)$



Task: given  $x$ , find  $y$

$$y = \arg \max_{y \in Y} P(y|x)$$

$$= \arg \max_{y \in Y} \frac{P(x, y)}{P(x)}$$

$$= \arg \max_{y \in Y} P(x, y)$$

# HMM – Viterbi Algorithm

$$\tilde{y} = \arg \max_{y \in \mathbb{Y}} P(x, y)$$

- Enumerate all possible  $y$ 
  - Assume there are  $|S|$  tags, and the length of sequence  $y$  is  $L$
  - There are  $|S|^L$  possible  $y$
- ***Viterbi algorithm***
  - Solve the above problem with complexity  $O(L|S|^2)$

# HMM - Summary

Problem 1:  
Evaluation



Problem 2:  
Inference



Problem 3:  
Training

$$F(x, y) = P(x, y) = P(y)P(x|y)$$

$$\tilde{y} = \arg \max_{y \in \mathbb{Y}} P(x, y)$$

$P(y)$  and  $P(x|y)$  can be simply  
obtained from training data



# HMM - Drawbacks

- Inference:

$$\tilde{y} = \arg \max_{y \in \mathbb{Y}} P(x, y)$$

- To obtain correct results ...

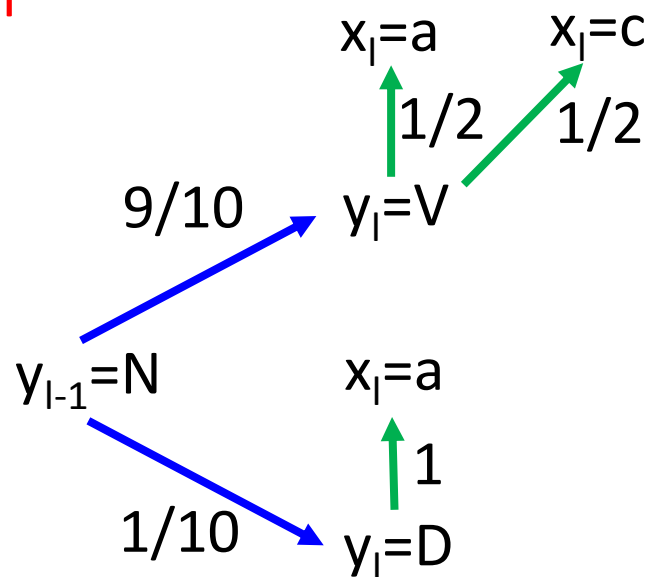
$(x, \hat{y}): P(x, \hat{y}) > \underline{P(x, y)}$  Can HMM guarantee that?  
not necessarily small

## Transition probability:

$P(V|N)=9/10$     $P(D|N)=1/10$  .....

## Emission probability:

$P(a|V)=1/2$     $P(a|D)=1$  .....



# HMM - Drawbacks

- Inference:

$$\tilde{y} = \arg \max_{y \in \mathbb{Y}} P(x, y)$$

- To obtain correct results ...

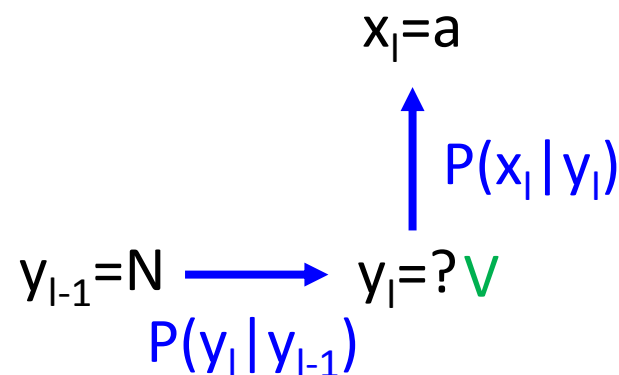
$(x, \hat{y}): P(x, \hat{y}) > \underline{P(x, y)}$  Can HMM guarantee that?  
not necessarily small

## Transition probability:

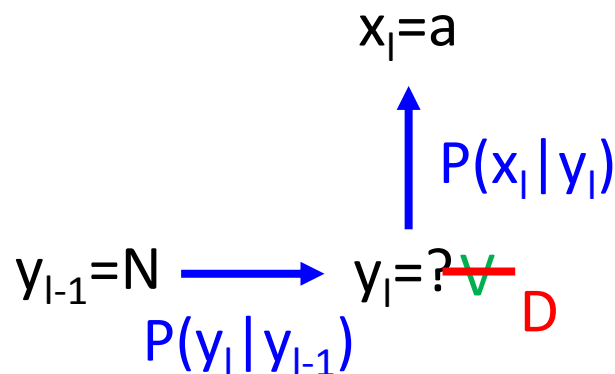
$$P(V|N)=9/10 \quad P(D|N)=1/10 \quad \dots\dots$$

## Emission probability:

$$P(a|V)=1/2 \quad P(a|D)=1 \quad \dots\dots$$



# HMM - Drawbacks



- Inference:

$$\tilde{y} = \arg \max_{y \in \mathbb{Y}} P(x, y)$$

- To obtain correct results ...

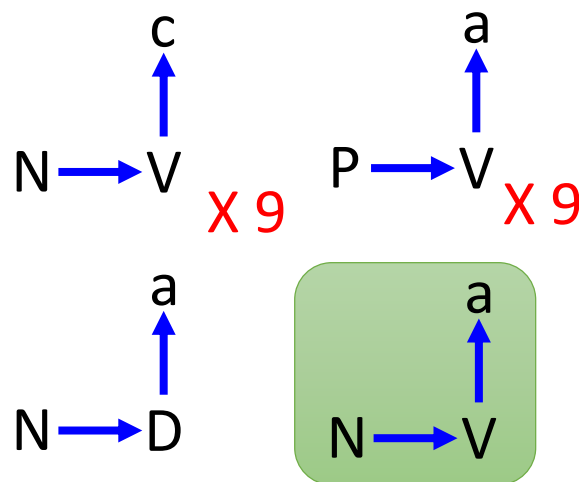
$(x, \hat{y}): P(x, \hat{y}) > \underline{P(x, y)}$  Can HMM guarantee that?  
 not necessarily small

**Transition probability:**

$$P(V|N)=9/10 \quad P(D|N)=1/10 \quad \dots\dots$$

**Emission probability:**

$$P(a|V)=1/2 \quad P(a|D)=1 \quad \dots\dots$$

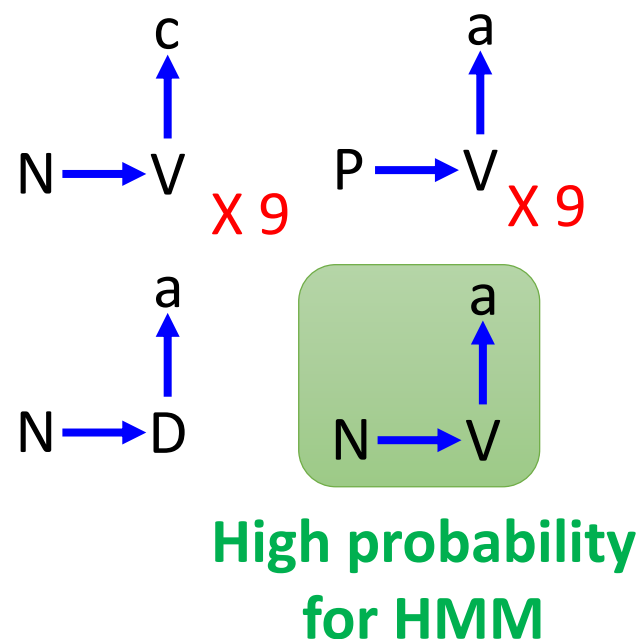


High probability  
for HMM

# HMM - Drawbacks

- The  $(x,y)$  never seen in the training data can have large probability  $P(x,y)$ .
- Benefit:
  - When there is only little training data

- More complex model can deal with this problem
- However, CRF can deal with this problem based on the same model



# Outline

Hidden Markov Model (HMM)



Conditional Random Field (CRF)



Structured Perceptron/SVM



Towards Deep Learning

# CRF

$$P(x, y) \propto \exp(\mathbf{w} \cdot \phi(x, y))$$

➤  $\phi(x, y)$  is a feature vector. What does it look like?

➤  $\mathbf{w}$  is a weight vector to be learned from training data

➤  $\exp(\mathbf{w} \cdot \phi(x, y))$  is always positive, can be larger than 1


$$\begin{aligned} P(y|x) &= \frac{P(x, y)}{\sum_{y'} P(x, y')} & P(x, y) &= \frac{\exp(\mathbf{w} \cdot \phi(x, y))}{R} \\ &= \frac{\exp(\mathbf{w} \cdot \phi(x, y))}{\sum_{y' \in \mathbb{Y}} \exp(\mathbf{w} \cdot \phi(x, y'))} & &= \frac{\exp(\mathbf{w} \cdot \phi(x, y))}{Z(x)} \end{aligned}$$

# $P(x, y)$ for CRF

$$P(x, y) \propto \exp(\mathbf{w} \cdot \boldsymbol{\phi}(x, y))$$

very different from HMM?

In HMM:


$$P(x, y) = P(y_1 | start) \prod_{l=1}^{L-1} P(y_{l+1} | y_l) P(end | y_L) \prod_{l=1}^L P(x_l | y_l)$$
$$\log P(x, y)$$

$$= \log P(y_1 | start) + \sum_{l=1}^{L-1} \log P(y_{l+1} | y_l) + \log P(end | y_L)$$
$$+ \sum_{l=1}^L \log P(x_l | y_l)$$

# $P(x, y)$ for CRF

$$\log P(x, y) = \log P(y_1 | \text{start}) + \sum_{l=1}^{L-1} \log P(y_{l+1} | y_l) + \log P(\text{end} | y_L) + \sum_{l=1}^L \log P(x_l | y_l)$$

Log probability of  
word  $t$  given tag  $s$

Number of tag  $s$  and word  $t$   
appears together in  $(x, y)$

$$\sum_{l=1}^L \log P(x_l | y_l) = \sum_{s, t} \log P(t | s) \times N_{s, t}(x, y)$$

Enumerate all possible tags  $s$   
and all possible word  $t$



# $P(x,y)$ for CRF

x: The dog ate the homework.  
      ↓     ↓     ↓     ↓     ↓  
 y: D   N   V   D   N

$$N_{D,the}(x,y) = 2$$

$$N_{N,dog}(x,y) = 1$$

$$N_{V,ate}(x,y) = 1$$

$$N_{N,homework}(x,y) = 1$$

$$N_{s,t}(x,y) = 0$$

(for any other s and t)

$$\sum_{l=1}^L \log P(x_l | y_l)$$

$$= \log P(\text{the} | D) + \log P(\text{dog} | N) + \log P(\text{ate} | V) \\ + \log P(\text{the} | D) + \log P(\text{homework} | N)$$

$$= \log P(\text{the} | D) \times 2 + \log P(\text{dog} | N) \times 1 + \log P(\text{ate} | V) \times 1 \\ + \log P(\text{homework} | N) \times 1$$

$$= \sum_{s,t} \log P(t | s) \times N_{s,t}(x,y)$$

# $P(x, y)$ for CRF

$$\log P(x, y)$$

$$= \boxed{\log P(y_1 | \text{start})} + \boxed{\sum_{l=1}^{L-1} \log P(y_{l+1} | y_l)} + \boxed{\log P(\text{end} | y_L)} \\ + \sum_{l=1}^L \log P(x_l | y_l)$$

$$\log P(y_1 | \text{start}) = \sum_s \log P(s | \text{start}) \times N_{\text{start}, s}(x, y)$$

$$\sum_{l=1}^{L-1} \log P(y_{l+1} | y_l) = \sum_{s, s'} \log P(s' | s) \times N_{s, s'}(x, y)$$

$$\log P(\text{end} | y_L) = \sum_s \log P(\text{end} | s) \times N_{s, \text{end}}(x, y)$$

# P(x,y) for CRF

$$\begin{aligned}
 \log P(x, y) &= \sum_{s,t} \log P(t|s) \times N_{s,t}(x, y) \\
 &+ \sum_s \log P(s|start) \times N_{start,s}(x, y) \\
 &+ \sum_{s,s'} \log P(s'|s) \times N_{s,s'}(x, y) \\
 &+ \sum_s \log P(end|s) \times N_{s,end}(x, y)
 \end{aligned}
 = \begin{bmatrix} \vdots \\ \log P(t|s) \\ \vdots \\ \vdots \\ \log P(s|start) \\ \vdots \\ \vdots \\ \log P(s'|s) \\ \vdots \\ \vdots \\ \log P(end|s) \\ \vdots \end{bmatrix} \cdot \begin{bmatrix} \vdots \\ N_{s,t}(x, y) \\ \vdots \\ \vdots \\ N_{start,s}(x, y) \\ \vdots \\ \vdots \\ N_{s,s'}(x, y) \\ \vdots \\ \vdots \\ N_{s,end}(x, y) \\ \vdots \end{bmatrix}$$

$$= w \cdot \phi(x, y)$$

$$P(x, y) = \exp(w \cdot \phi(x, y))$$

# $P(x, y)$ for CRF

$\propto$

$$P(x, y) \propto \exp(w \cdot \phi(x, y))$$

However, we do not give  $w$  any constraints during training

$$\phi(x, y) = \begin{bmatrix} \vdots \\ N_{s,t}(x, y) \\ \vdots \\ N_{start,s}(x, y) \\ \vdots \\ N_{s,s'}(x, y) \\ \vdots \\ N_{s,end}(x, y) \\ \vdots \end{bmatrix}$$

$$w = \begin{bmatrix} \vdots \\ w_{s,t} \\ \vdots \\ w_{start,s} \\ \vdots \\ w_{s,s'} \\ \vdots \\ w_{s,end} \\ \vdots \end{bmatrix}$$

$$\longrightarrow \log P(x_i = t | y_i = s)$$

$$P(x_i = t | y_i = s) = e^{w_{s,t}} \text{ means}$$

$$\longrightarrow \log P(s | start)$$

$$P(s | start) = e^{w_{start,s}} \text{ means}$$

$$\longrightarrow \log P(y_i = s' | y_{i-1} = s)$$

$$P(y_i = s' | y_{i-1} = s) = e^{w_{s,s'}} \text{ means}$$

$$\longrightarrow \log P(end | s)$$

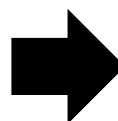
... ..

# Feature Vector

- What does  $\phi(x, y)$  look like?

x: The dog ate the homework.  
↓ ↓ ↓ ↓ ↓  
y: D N V D N

- $\phi(x, y)$  has two parts
    - Part 1: relations between tags and words
    - Part 2: relations between tags
- If there are  $|S|$  possible tags,  
 $|L|$  possible words  
Part 1 has  $|S| \times |L|$  dimensions



Part 1	Value
D, the	2
D, dog	0
D, ate	0
D, homework	0
.....	.....
N, the	0
N, dog	1
N, ate	0
N, homework	1
.....	.....
V, the	0
V, dog	0
V, ate	1
V, homework	0
.....	.....

# Feature Vector

- What does  $\phi(x, y)$  look like?

x: The dog ate the homework.  
      ↓     ↓     ↓     ↓     ↓  
 y: D     N     V     D     N

- $\phi(x, y)$  has two parts
  - Part 1: relations between tags and words
  - Part 2: relations between tags

$N_{s,s'}(x, y)$ : Number of tags  $s$  and  $s'$  consecutively in  $(x, y)$

$N_{D,D}(x, y) \rightarrow$

$N_{D,N}(x, y) \rightarrow$

Part 2	Value
D, D	0
D, N	2
D, V	0
.....	.....
N, D	0
N, N	0
N, V	1
.....	.....
V, D	1
V, N	0
V, V	0
.....	.....
Start, D	1
Start, N	0
.....	.....
End, D	0
End, N	1

# Feature Vector

- What does  $\phi(x, y)$  look like?

x: The dog ate the homework.  
      ↓     ↓     ↓     ↓     ↓  
 y: D     N     V     D     N

- $\phi(x, y)$  has two parts
  - Part 1: relations between tags and words
  - Part 2: relations between tags

If there are  $|S|$  possible tags,  
 $|S| \times |S| + 2 \times |S|$  dimensions

Define any  $\phi(x, y)$  you like!

Part 2	Value
D, D	0
D, N	2
D, V	0
.....	.....
N, D	0
N, N	0
N, V	1
.....	.....
V, D	1
V, N	0
V, V	0
.....	.....
Start, D	1
Start, N	0
.....	.....
End, D	0
End, N	1

$$P(y|x)$$

$$= \frac{P(x, y)}{\sum_{y'} P(x, y')}$$

# CRF – Training Criterion

- Given training data:  $\{(x^1, \hat{y}^1), (x^2, \hat{y}^2), \dots (x^N, \hat{y}^N)\}$
- Find the weight vector  $w^*$  maximizing objective function  $O(w)$ :

$$w^* = \arg \max_w O(w) \quad O(w) = \sum_{n=1}^N \log P(\hat{y}^n | x^n)$$

$$\log P(\hat{y}^n | x^n) = \log \underbrace{P(x^n, \hat{y}^n)}_{\text{Maximize what we observe}} - \log \underbrace{\sum_{y'} P(x^n, y')}_{\text{Minimize what we don't observe}}$$

Maximize what  
we observe

Minimize what we  
don't observe



# CRF – Gradient Ascent

## Gradient descent

Find a set of parameters  $\theta$  minimizing cost function  $C(\theta)$

$$\theta \rightarrow \theta - \eta \nabla C(\theta)$$

Opposite direction of the gradient

## Gradient Ascent

Find a set of parameters  $\theta$  maximizing objective function  $O(\theta)$

$$\theta \rightarrow \theta + \eta \nabla O(\theta)$$

The same direction of the gradient

# CRF - Training

$$O(w) = \sum_{n=1}^N \log P(\hat{y}^n | x^n) = \sum_{n=1}^N O^n(w)$$

Compute

$$\nabla O^n(w) = \begin{bmatrix} \vdots \\ \partial O^n(w) / \partial w_{s,t} \\ \vdots \\ \partial O^n(w) / \partial w_{s,s'} \\ \vdots \end{bmatrix}$$

Let me show  $\frac{\partial O^n(w)}{\partial w_{s,t}}$

$\frac{\partial O^n(w)}{\partial w_{s,s'}}$  very similar

# CRF - Training

$$P(y'|x^n) = \frac{\exp(w \cdot \phi(x^n, y'))}{Z(x^n)}$$

$$w_{s,t} \rightarrow w_{s,t} + \eta \frac{\partial O(w)}{\partial w_{s,t}}$$

After some math .....

Can be computed by Viterbi algorithm as well

$$\frac{\partial O^n(w)}{\partial w_{s,t}} = \underline{N_{s,t}(x^n, \hat{y}^n)} - \sum_{y'} \underline{P(y'|x^n) N_{s,t}(x^n, y')}$$

If word  $t$  is labeled by tag  $s$  in training examples  $(x^n, \hat{y}^n)$ , then increase  $w_{s,t}$

If word  $t$  is labeled by tag  $s$  in  $(x^n, y')$  which not in training examples, then decrease  $w_{s,t}$

$$P(y'|x^n) = \frac{\exp(w \cdot \phi(x^n, y'))}{Z(x^n)}$$

## CRF - Training

$$\nabla O(w) = \phi(x^n, \hat{y}^n) - \sum_{y'} P(y'|x^n) \phi(x^n, y')$$

### Stochastic Gradient Ascent

Random pick a data  $(x^n, \hat{y}^n)$

$$w \rightarrow w + \eta \left( \phi(x^n, \hat{y}^n) - \sum_{y'} P(y'|x^n) \phi(x^n, y') \right)$$

# CRF – Inference

- Inference

$$y = \arg \max_{y \in Y} P(y|x) = \arg \max_{y \in Y} P(x, y)$$

$$= \arg \max_{y \in Y} w \cdot \phi(x, y) \quad \text{Done by Viterbi as well}$$

$$P(x, y) \propto \exp(w \cdot \phi(x, y))$$

# CRF v.s. HMM

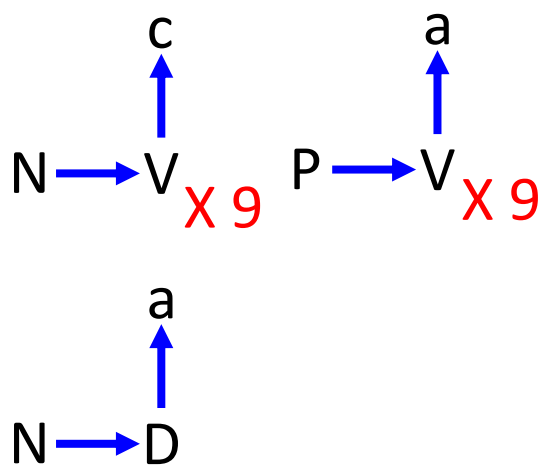
- CRF: increase  $P(x, \hat{y})$ , decrease  $P(x, y')$

HMM does not do that

- To obtain correct results ...

$$(x, \hat{y}): P(x, \hat{y}) > P(x, y)$$

CRF more likely to achieve that than HMM



HMM:

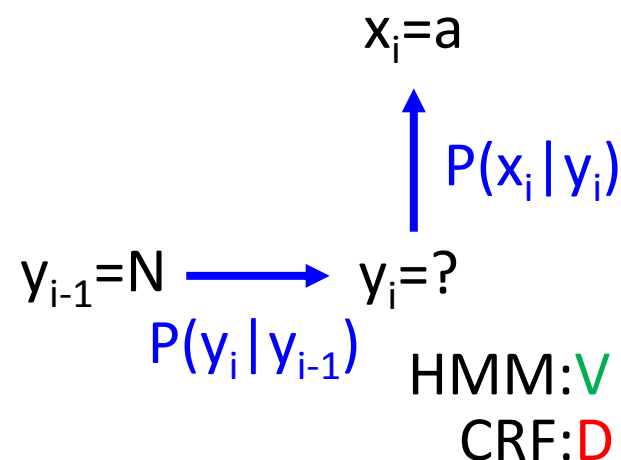
$P(V | N) = 9/10$

$P(D | N) = 1/10$

$P(a | V) = 1/2 \rightarrow 0.1$

$P(a | D) = 1$

CRF:



HMM: **V**

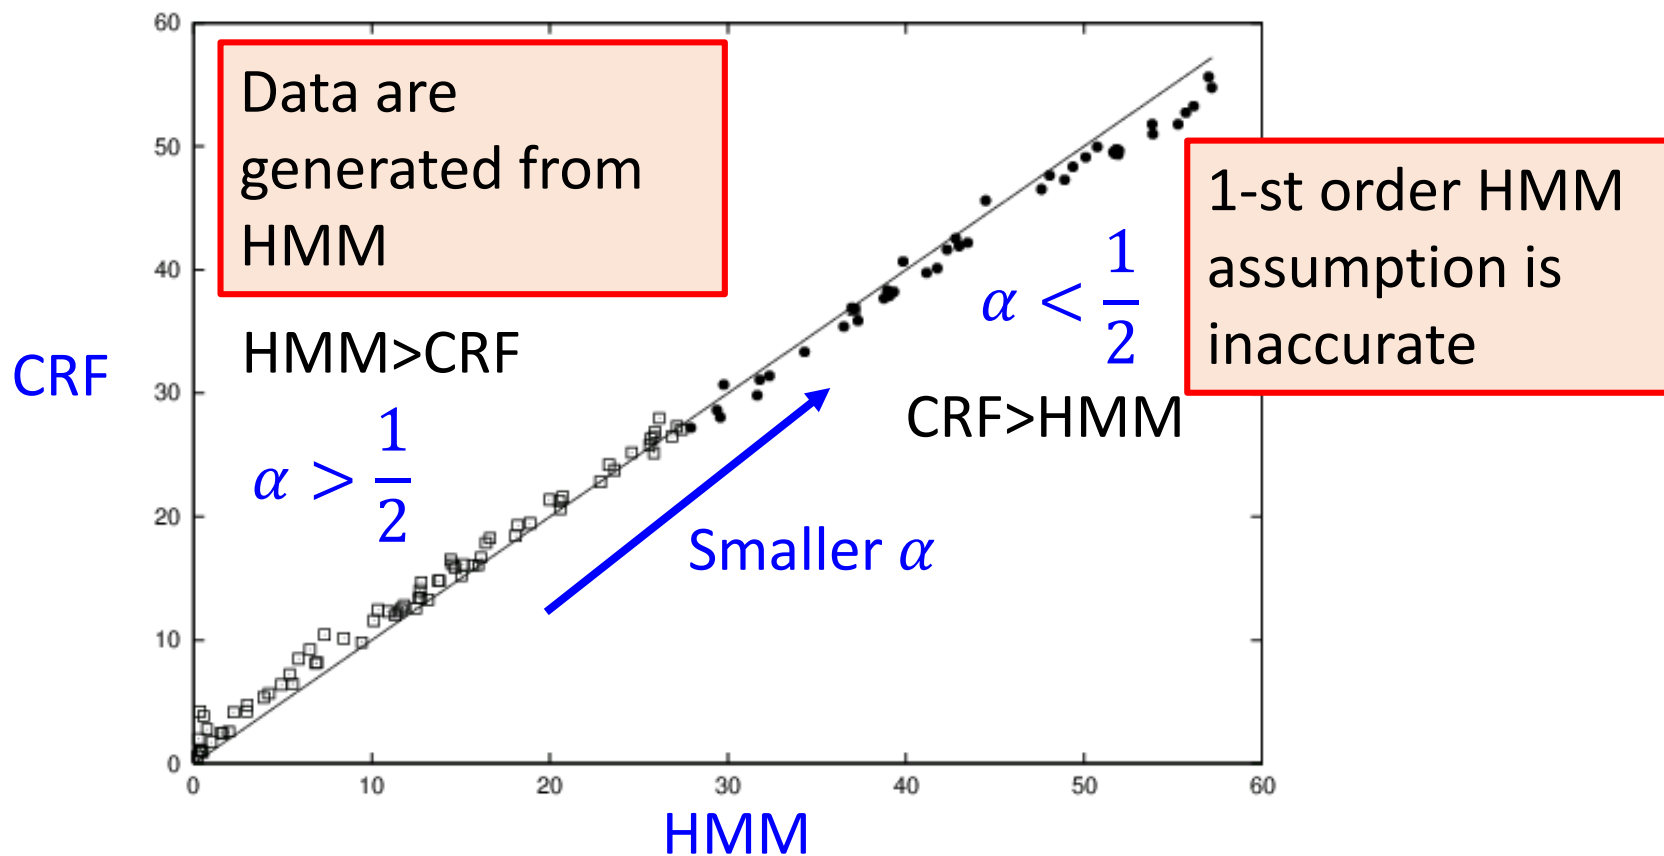
CRF: **D**

# Synthetic Data

- $x_i \in \{a - z\}, y_i \in \{A - E\}$
- Generating data from a mixed-order HMM
  - Transition probability:
    - $\alpha P(y_i | y_{i-1}) + (1 - \alpha) P(y_i | y_{i-1}, y_{i-2})$
  - Emission probability:
    - $\alpha P(x_i | y_i) + (1 - \alpha) P(x_i | y_i, x_{i-1})$
- Comparing HMM and CRF
  - All the approaches only consider 1-st order information
    - Only considering the relation of  $y_{i-1}$  and  $y_i$
  - In general, all the approaches have worse performance with smaller  $\alpha$

Ref: John D. Lafferty, Andrew McCallum, and Fernando C. N. Pereira, “*Conditional Random Fields: Probabilistic Models for Segmenting and Labeling Sequence Data*”, ICML, 2001

# Synthetic Data: CRF v.s. HMM





# CRF - Summary

Problem 1:  
Evaluation

$$F(x, y) = P(y|x) = \frac{\exp(w \cdot \phi(x, y))}{\sum_{y' \in \mathbb{Y}} \exp(w \cdot \phi(x, y'))}$$

Problem 2:  
Inference

$$\tilde{y} = \arg \max_{y \in \mathbb{Y}} P(y|x) = \arg \max_{y \in \mathbb{Y}} w \cdot \phi(x, y)$$

Problem 3:  
Training

$$w^* = \arg \max_w \prod_{n=1}^N P(\hat{y}^n | x^n)$$

$$w \rightarrow w + \eta \left( \phi(x^n, \hat{y}^n) - \sum_{y'} P(y'|x^n) \phi(x^n, y') \right)$$

# Outline

Hidden Markov Model (HMM)



Conditional Random Field (CRF)



Structured Perceptron/SVM



Towards Deep Learning

# Structured Perceptron

Problem 1:  
Evaluation

$$F(x, y) = w \cdot \underline{\phi(x, y)}$$

The same as CRF

Problem 2:  
Inference

$$\tilde{y} = \arg \max_{y \in \mathbb{Y}} w \cdot \phi(x, y)$$

Viterbi

Problem 3:  
Training

$$\forall n, \forall y \in \mathbb{Y}, y \neq \hat{y}^n:$$

$$w \cdot \phi(x^n, \hat{y}^n) > w \cdot \phi(x^n, y)$$

$$\tilde{y}^n = \arg \max_y w \cdot \phi(x^n, y)$$

$$w \rightarrow w + \phi(x^n, \hat{y}^n) - \phi(x^n, \tilde{y}^n)$$

# Structured Perceptron v.s. CRF

- Structured Perceptron

$$\tilde{y}^n = \arg \max_y w \cdot \phi(x^n, y)$$

$$w \rightarrow w + \underbrace{\phi(x^n, \hat{y}^n)}_{\text{Hard}} - \underbrace{\phi(x^n, \tilde{y}^n)}_{\text{Hard}}$$

Hard

- CRF

$$w \rightarrow w + \eta \left( \underbrace{\phi(x^n, \hat{y}^n)}_{\text{Soft}} - \underbrace{\sum_{y'} P(y'|x^n) \phi(x^n, y')}_{\text{Soft}} \right)$$

Soft

# Structured SVM

Problem 1:  
Evaluation

$$F(x, y) = w \cdot \underline{\phi(x, y)}$$

The same as CRF

Problem 2:  
Inference

$$\tilde{y} = \arg \max_{y \in \mathbb{Y}} w \cdot \phi(x, y)$$

Viterbi

Problem 3:  
Training

Consider margin and error:

Way 1. Gradient Descent

Way 2. Quadratic Programming  
(Cutting Plane Algorithm)

# Structured SVM – Error Function

- Error function:  $\Delta(\hat{y}^n, y)$ 
  - $\Delta(\hat{y}^n, y)$ : Difference between  $y$  and  $\hat{y}^n$
  - Cost function of structured SVM is the upper bound of  $\Delta(\hat{y}^n, y)$
  - Theoretically,  $\Delta(y, \hat{y}^n)$  can be any function you like
  - However, you need to solve **Problem 2.1**
    - $\bar{y}^n = \underset{y}{\operatorname{argmax}} [\Delta(\hat{y}^n, y) + w \cdot \phi(x^n, y)]$

## Example

$\hat{y}$ : A T T C G G G G A T

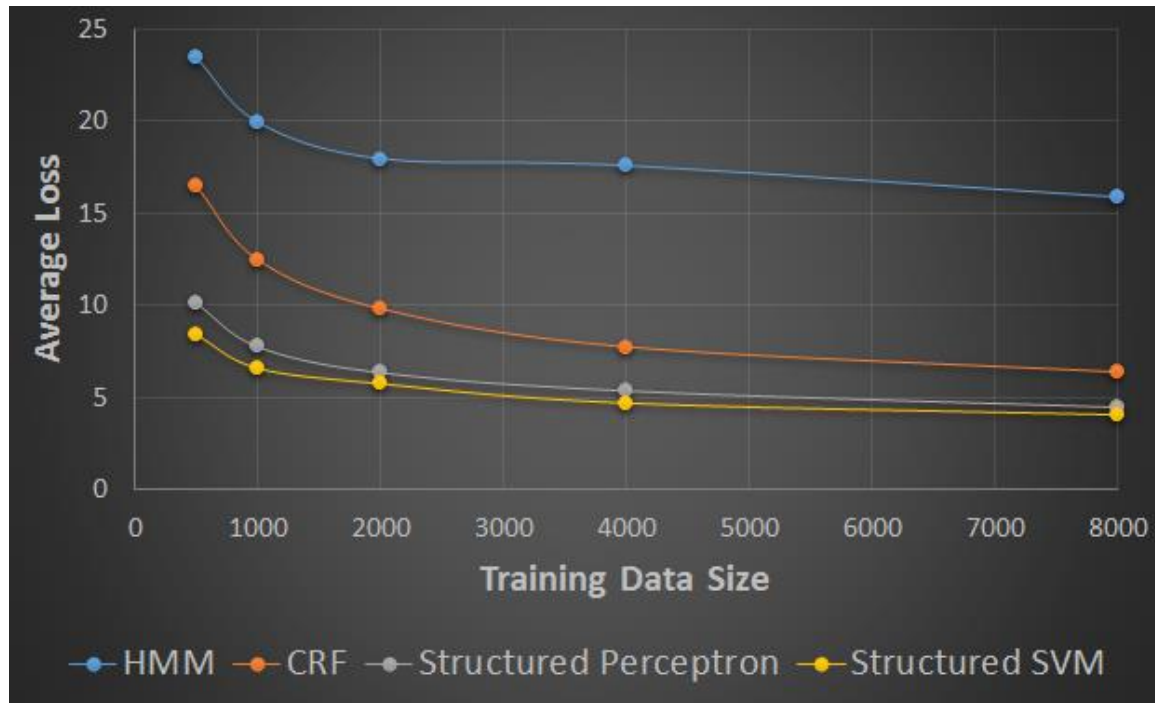
$\Delta(\hat{y}, y) = 3/10$       $y$ : A T T A G G A G A A

In this case, problem 2.1 can be solved by Viterbi Algorithm

# Performance of Different Approaches

## POS Tagging

Ref: Nguyen, Nam, and Yunsong Guo.  
"Comparisons of sequence labeling  
algorithms and extensions." *ICML*, 2007.



## Name Entity Recognition

Method	HMM	CRF	Perceptron	SVM
Error	9.36	5.17	5.94	5.08

Ref: Tsochantaridis, Ioannis, et al. "Large margin methods for structured and interdependent output variables." *Journal of Machine Learning Research*. 2005.

# Outline

Hidden Markov Model (HMM)



Conditional Random Field (CRF)



Structured Perceptron/SVM



Towards Deep Learning



# How about RNN?

- RNN, LSTM

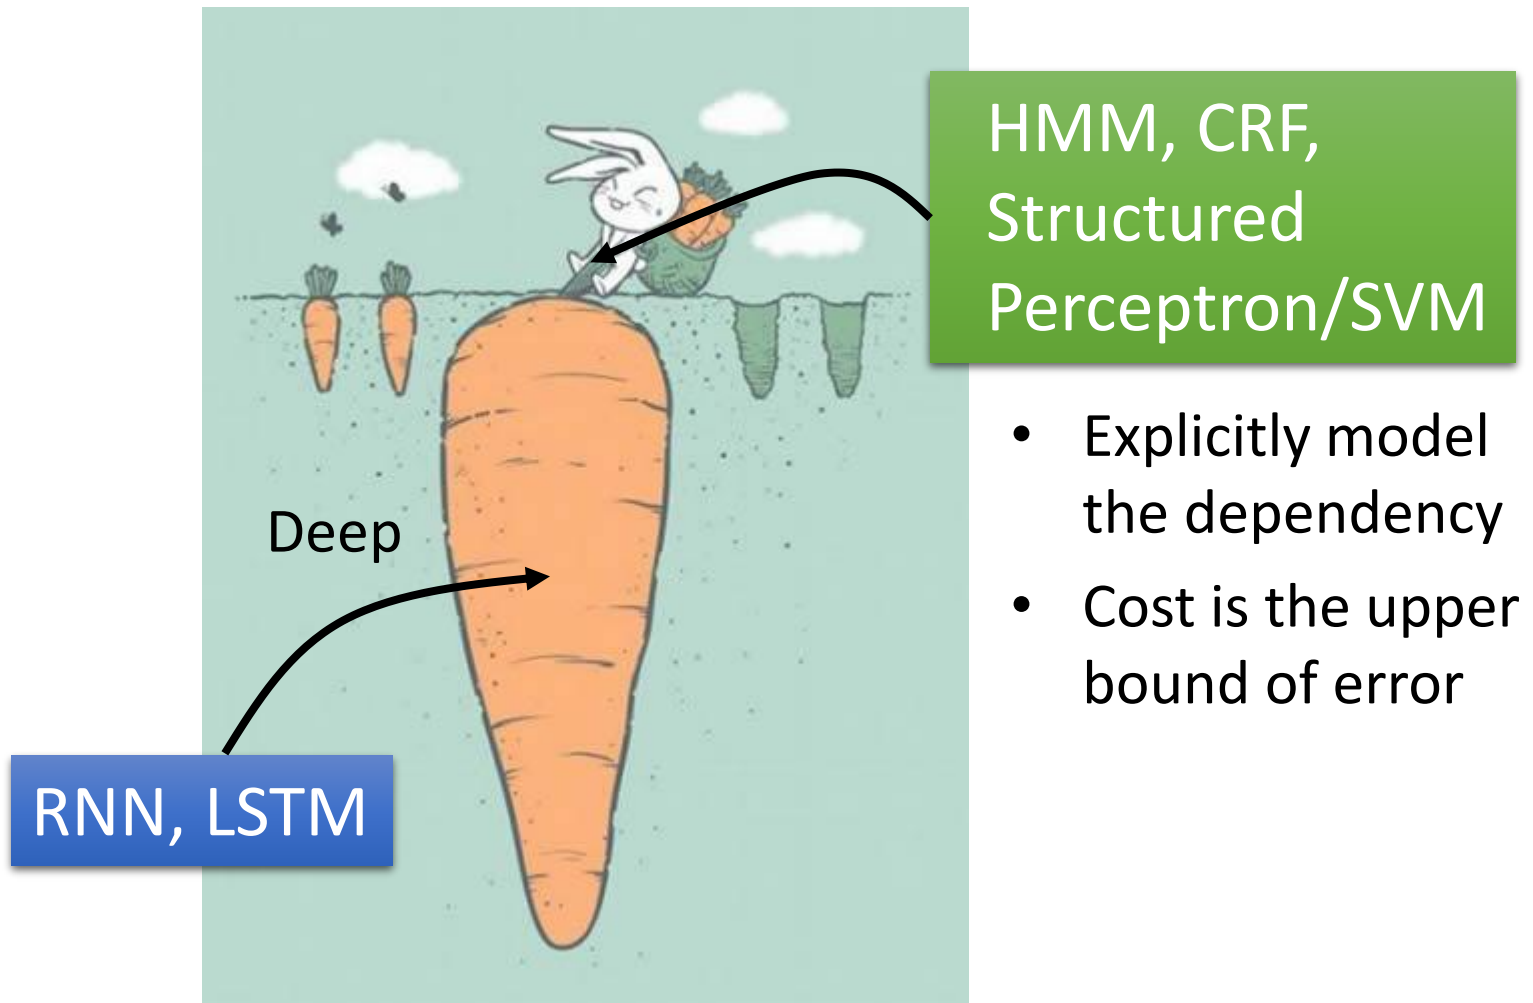
- Unidirectional RNN does not consider the whole sequence
- Cost and error not always related
- Deep 勝



- HMM, CRF, Structured Perceptron/SVM

- Using Viterbi, so consider the whole sequence 勝?
- How about Bidirectional RNN?
- Can explicitly consider the label dependency 勝
- Cost is the upper bound of error 勝

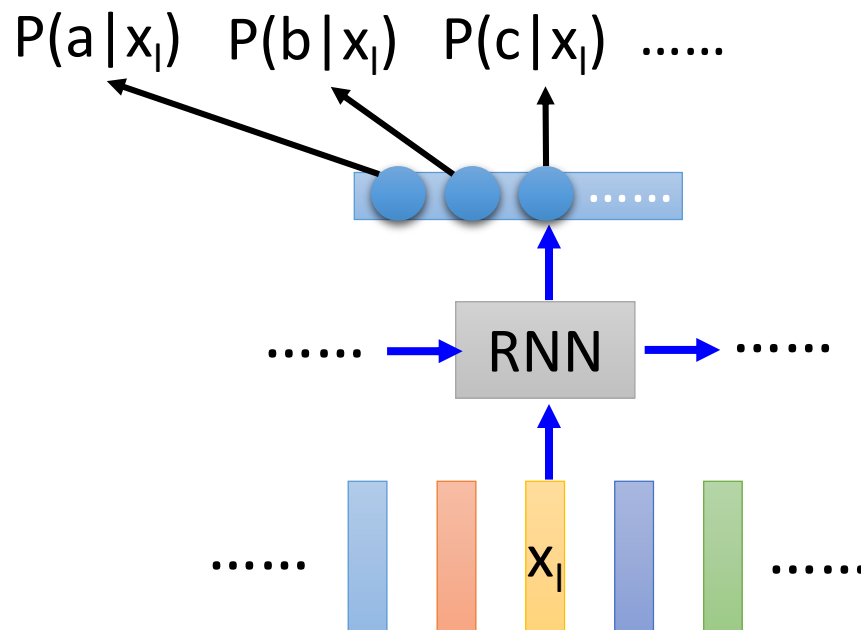
# Integrated together



# Integrated together

- Speech Recognition: CNN/RNN or LSTM/DNN + HMM

$$P(x, y) = P(y_1 | start) \prod_{l=1}^{L-1} P(y_{l+1} | y_l) P(end | y_L) \prod_{l=1}^L \underline{P(x_l | y_l)}$$

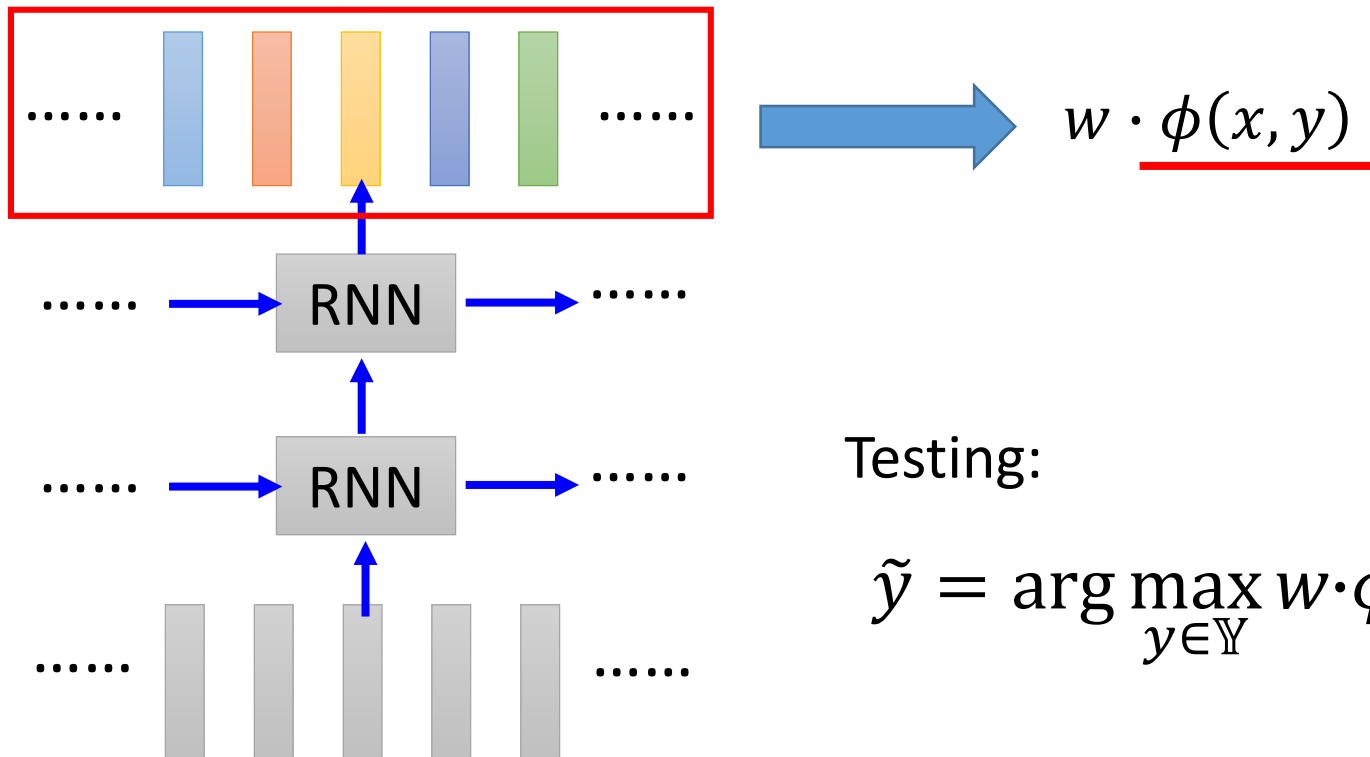


$$P(x_l | y_l) = \frac{P(x_l, y_l)}{P(y_l)}$$

$$= \frac{\text{RNN} \quad P(y_l | x_l) \cancel{P(x_l)}}{\text{Count} \quad P(y_l)}$$

# Integrated together

- Semantic Tagging: Bi-directional RNN/LSTM + CRF/Structured SVM



# Concluding Remarks

	Problem 1	Problem 2	Problem 3
HMM	$F(x, y) = P(x, y)$	Viterbi	Just count
CRF	$F(x, y) = P(y x)$	Viterbi	Maximize $P(\hat{y} x)$
Structured Perceptron	$F(x, y) = w \cdot \phi(x, y)$ (not a probability)	Viterbi	$F(x, \hat{y}) > F(x, y')$
Structured SVM	$F(x, y) = w \cdot \phi(x, y)$ (not a probability)	Viterbi	$F(x, \hat{y}) > F(x, y')$ with <b>margins</b>

The above approaches can combine with deep learning to have better performance.

# Acknowledgement

- 感謝 曹燁文 同學於上課時發現投影片上的錯誤
- 感謝 **Ryan Sun** 來信指出投影片上的錯誤

# Appendix

# CRF - Training

$$O^n(w) = \log \frac{\exp(w \cdot \phi(x^n, \hat{y}^n))}{Z(x^n)} \quad Z(x^n) = \sum_{y'} \exp(w \cdot \phi(x^n, y'))$$
$$= \underline{w \cdot \phi(x^n, \hat{y}^n)} - \log Z(x^n)$$

$$\frac{\partial O^n(w)}{\partial w_{s,t}} = \underline{N_{s,t}(x^n, \hat{y}^n)}$$



The number of word t labeled as s in  $(x^n, \hat{y}^n)$

The value of the dimension in  $\phi(x^n, \hat{y}^n)$  corresponding to  $w_{s,t}$ .

$$w \cdot \phi(x^n, \hat{y}^n) = \sum_{s,t} w_{s,t} \cdot N_{s,t}(x^n, \hat{y}^n) + \sum_{s,s'} w_{s,s'} \cdot N_{s,s'}(x^n, \hat{y}^n)$$



# CRF - Training

$$O^n(w) = \log \frac{\exp(w \cdot \phi(x^n, \hat{y}^n))}{Z(x^n)} \quad Z(x^n) = \sum_{y'} \exp(w \cdot \phi(x^n, y'))$$

$$= \underline{w \cdot \phi(x^n, \hat{y}^n)} - \underline{\log Z(x^n)}$$

$$\frac{\partial O^n(w)}{\partial w_{s,t}} = \underline{N_{s,t}(x^n, \hat{y}^n)} - \frac{1}{\underline{Z(x^n)}} \underline{\frac{\partial Z(x^n)}{\partial w_{s,t}}}$$

$$= \sum_{y'} \left[ \frac{\exp(w \cdot \phi(x^n, y'))}{Z(x^n)} \right] N_{s,t}(x^n, y') = \sum_{y'} \left[ P(y'|x^n) \right] N_{s,t}(x^n, y')$$

$$\underline{\frac{\partial Z(x^n)}{\partial w_{s,t}}} = \sum_{y'} \exp(w \cdot \phi(x^n, y')) N_{s,t}(x^n, y')$$

# CRF v.s. HMM

- Define  $\phi(x, y)$  you like
  - For example, besides the features just described, there are some useful extra features in POS tagging.
    - Number of times a capitalized word is labeled as Noun
    - Number of times a word end with **ing** is labeled as Noun
- Can you consider this kind of features by HMM? Too sparse...

$$P(x_i = A, x_i \text{ is capitalized}, x_i \text{ end with } \mathbf{ing}, \dots | y_i = N)$$

## Method 1:

$$P(x_i = A | y_i = N) P(x_i \text{ is capitalized} | y_i = N) \dots$$

Inaccurate assumption

Method 2. Give the distribution some assumptions?